

RAG System Workflow

A RAG system workflow begins with document ingestion and preprocessing. Documents are split into chunks and converted into embeddings. These embeddings are stored in a vector database. When a user submits a query, relevant chunks are retrieved. The language model then generates an answer using the retrieved context.